



USER GUIDE AND HANDOVER

Collection Manager

A receivables system built on the daily outstanding export from Marg ERP. It records the credit terms Marg cannot hold, measures every bill against the term that actually applies to that party, and refuses to call anything overdue until someone has told it what the term is.

PARTIES	OPEN BILLS	NET OUTSTANDING	RUNNING COST
819 709 owing · 97 in credit	8,197 back to Mar 2020	₹7.93 Cr matches Marg exactly	₹0 per month

Why this exists

Marg tells you what is owed. It cannot tell you what is **late**, because it holds no credit terms — its Due Date column simply repeats the Bill Date on almost every row. So the only thing the export can honestly say is how *old* a bill is.

Old and late are not the same thing. A hospital that submits bills by the 5th and is paid on the 25th of the following month will routinely carry a 78-day-old bill that is perfectly on time. Chasing them for it wastes the call and damages the relationship.

This system keeps the two ideas apart. Everything follows from that one distinction, including the places where it refuses to answer.

What the file itself revealed

Two things in the real export are worth knowing, because both would quietly produce wrong figures if handled carelessly.

The party totals and the bill rows disagree. For 46 parties, Marg's own header total does not match the sum of the bills beneath it — by about ₹1 crore in aggregate. Adding up bill rows gives ₹8.93 Cr where Marg shows ₹7.93 Cr. This system takes the header figure, as Marg does, and lists every disagreement so nothing is hidden.

Marg already holds a due date for 19 parties, implying credit periods of 60 or 30 days. Those are shown for review but deliberately not applied — a term becomes real only when a person approves it.

Getting started

- ① **Open the site and sign in.** Each person gets their own login. What they can change depends on their role.
- ② **Fill in credit terms.** This is the one thing only Neomed can supply, and until it is done the system can show balances but not lateness. Go to **Credit master** — parties are listed largest first, so the work is finite: the top hundred rows cover most of the money.
- ③ **Import the file each morning.** Export the outstanding bill-wise report from Marg, drop it into **Import**, read the preview, confirm. About a minute.
- ④ **Work the call list.** **Today** shows who to call and why. Log what happened as you go.

The system gets better every day you use it. Marg exports no receipts, so payment history is worked out by comparing today's file against yesterday's — a balance that fell is money that arrived. One file tells it nothing about behaviour. A month of files tells it a great deal.

The screens

Today

The current position, and what to do about it. Four headline figures, the portfolio ageing strip, the priority call list with a written reason for each party, and where the money is concentrated.

The **Work this first** list is ranked by money at risk — the amount, weighted by how far past term it is, weighted again by how that party usually behaves. Each row explains itself in a sentence built from that party's own figures.

DELIBERATELY LEFT OFF THE CALL LIST

- **Held by an open claim.** The next action is ours, not theirs. They are not defaulters.
- **In dispute, legal, or written off.** Handled elsewhere.
- **No credit term recorded.** Nothing can be called late without one.
- **Below the small-balance threshold.** Chasing costs more than the money.
- **In credit or at zero.** They owe nothing.

Parties

All 819, searchable and sortable on any column. Filters for the slices you actually want: critical, above ₹5 L, term not set, credit balance. Clicking a row opens a drawer over the list rather than navigating away, so you keep your place, your filter and your scroll position.

The drawer shows the ageing split, every open bill, and — where the header and bill rows disagree — both figures side by side with a note saying which one is being used. It also carries:

- **Contact** — the mobile number and contact person. Marg exports neither, so they are typed here once. An import can never overwrite them. **Call** opens the dialer; **WhatsApp** opens a draft naming the party and the balance — nothing is sent until you press send.
- **Activity** — one timeline of what the team did and what the money did: follow-ups, promises, claims, and payments derived from comparing exports.
- **Log follow-up** — record a call without leaving the party.
- **Ask the assistant** — hands that party's name straight to the assistant and asks for its payment history.

Credit master

The data-entry screen, built for speed. Two models, because retailers and hospitals are paid differently:

MODEL	HOW IT WORKS	WHO USES IT
Days	So many days from the bill date. 30, 45, 60, 90.	Retailers, chemists
Monthly cycle	Bills must reach them by a set day to make that month's run; payment is released on a set day, that month or the next.	Hospitals, institutions
Not set	No term recorded. The party cannot be judged.	The starting state

Select several rows and apply one term to all of them — most hospitals share a cycle. The progress header shows how much of the book becomes judgeable as you work, so the effort has a visible end.

Every party that owes money is here, not a fixed top slice. The list starts at 150 rows and grows on request, because several hundred rows of live input fields at once makes typing sluggish. Search by name, or tick *Only those without a term* to hide what is already done.

Import

Drag the file in. It is read in your browser and you see exactly what it contains before anything is saved: the totals, what changed since the previous file, and every discrepancy found. **Nothing is written until you press confirm**, so a bad export can be discarded without leaving a trace.

Follow-ups, Promises, Claims

Follow-ups lists what was missed first, then what is planned, then what the team has been doing. **Claims** groups by who owes the internal action, because that is what unblocks the money.

Promises is the one worth explaining. A promise closes only when a payment supports it. Nobody can mark one kept by hand — press *Check against payments* and the system matches each promise against money that actually arrived, and records which payments justify the result.

Settings

Team and roles, the small-balance threshold, the category defaults, and the full import history with the warnings each file raised.

The assistant

Ask in plain English. Answers come back as figures, ageing strips and ranked lists — not paragraphs.

THINGS IT ANSWERS

Give me today's calls on priority

The ranked call list, with a reason for each party.

Who has payment pending?

The same ranking, phrased the other way round.

Who did not keep their promise?

Broken and partly-kept promises, with what was actually received.

What is the payment history of Godavari?

How that party has paid over time: money received, bills cleared, median days to pay, and a rating.

How much does Godavari owe?

Today's balance and ageing — a different question from the one above, and answered differently.

Parties above 5 lakh

Understands "5 lakh", "₹5L", "50 lakh", "1 crore".

Parties with a bill older than 90 days

Bill age, clearly labelled as age rather than lateness.

Which parties cannot I judge?

Everything blocked for want of a credit term, largest first.

Who is blocked by a claim?

Money held by our own pending action.

What changed since the last import?

Money in, new billing, and the net movement.

Who is in credit?

The parties holding an advance rather than a debt.

Which accounts are too small to chase?

Below the threshold, with the cost argument.

How many parties are there?

Counts for the current snapshot.

Ask one thing at a time. *"Which party is our priority that has the highest outstanding"* is two different rankings in one sentence and will not match; ask for the call list or the top parties separately.

What it will not do

There is no language model behind this. Questions are matched against a fixed set of patterns and every figure is computed by the database. Three consequences, all deliberate:

- **The same question always gives the same answer.**
- **It cannot invent a number.** Every figure is read, never generated.
- **It says when it does not understand** rather than guessing. A confident answer to a question nobody asked is worse than admitting the miss.

Asked something it cannot honestly answer, it stops and says so.

Ask how overdue a party is when no term is recorded and it will tell you it cannot answer that, explain that bill age is not lateness, give you the figures it does have, and ask for the term. It will never fill the gap with an estimate. This is the single most important behaviour in the product.

Reading the ageing strip



Every bucket is measured against that party's **own expected due date**, never against raw bill age. The four segments always add up to the figure Marg shows, even where the bill rows disagree with the header.

A **hatched grey strip** means no credit term is recorded. It is not an empty strip and it is not zero — it means the question cannot be answered yet. Money in credit is marked in violet; it is an advance, not a debt.

The rules it will not break

RULE	
1	A credit term is never guessed. Not from bill dates, not from due dates, not from past behaviour. If it is not recorded, the system says so and asks.
2	A category default is not an approved term. Where one is used it is labelled an assumption everywhere it appears.
3	A party's total comes from Marg's own header figure , never from adding up bill rows.
4	Imported rows are never edited. Correcting a figure means importing a new file. Marg stays the financial truth.
5	A promise closes only when a payment supports it. Nobody can mark one kept by hand.
6	A credit balance is never counted as debt. 97 parties are in credit by ₹12.61 L between them.
7	Money reads the Indian way — ₹89.72 L, ₹7.93 Cr — and figures line up in columns.
8	Every figure traces to a snapshot and a rule. No estimates, and no rounding that hides a discrepancy.

Who can do what

ROLE	CAN CHANGE	TYPICALLY
Owner	Everything, including thresholds and who is on the team	The proprietor
Accounts	Credit terms, claims, follow-ups, promises, and imports	The accounts desk
Sales	Follow-ups and promises only	Field staff

Everyone sees the whole book. The limits are on *changing* things, and they are enforced in the database — a salesperson cannot set a credit term even if they reach the system directly. That is deliberate: the credit term is the one number the entire product depends on.

Adding someone

Create the user in Supabase under Authentication, then give them a role. A person with a login but no role can sign in and will see nothing — that is the security working, not a fault.

Adding someone to the team

1 **Create the login.** In Supabase, **Authentication → Users → Add user**. Enter their email and a password, and tick *Auto Confirm User* — without that they cannot sign in until they click a confirmation email.

2 **Give them a role.** In the SQL editor:

```
insert into app_users (id, full_name, role)
select id, 'S. Pawar', 'sales' from auth.users where email = 'sales@neomed.in';
```

Use `owner`, `accounts` or `sales`.

3 **Send them the link and their password.** They change it from the sign-in screen.

A person with a login but no role signs in and sees nothing at all. That is the security working, not a fault — the role is what grants access.

Where it stands today

PHASE	STATE	WHAT IT COVERS
Phase 1	Complete	Parser, import, database, ageing engine, parties list, party detail, credit master, dashboard
Phase 2	Complete	Claims, follow-ups, promises, roles and permissions
Phase 3	Built, waiting on data	Snapshot comparison, derived payment history, behaviour profiles, assistant

Two things are true on day one, and both are expected.

The book cannot be aged yet. No credit term has been recorded, so all ₹8.06 Cr sits under "cannot be judged". Typing in the terms for the largest parties is what turns this from a list of balances into a collection system.

Payment history is empty. It is derived by comparing consecutive exports, and only one file has been imported. Import tomorrow's and it begins; expect a few weeks before behaviour ratings mean much.

What it costs to run

ITEM	COST
Database, authentication, file storage	₹0
Hosting	₹0
Assistant — no language model, so nothing to meter	₹0
Domain, if wanted	~₹900/yr

Nothing bills per use. This dataset is about 25 MB a year against a 500 MB free allowance. The only reasons to ever pay would be outgrowing that, which takes years, or wanting daily backups.

Honest limitations

- **The export is manual.** If nobody exports the file, the book goes stale. Every screen shows the age of the snapshot it is reading for exactly that reason.
- **Payment history is inferred, not read.** Marg exports no receipts, so payments are derived from balances falling between two files. Reliable, but money paid and re-billed the same day is invisible.
- **Behaviour ratings need history.** A party stays unrated until three of its bills have been settled and observed.
- **Some bill series are unclassified.** 1,634 rows use voucher prefixes (CRA, CRB, CRC, CRD) that were not in the original specification. They are imported and counted correctly, but grouped as "other" until someone at Neomed says what those series mean.

Questions we need answered

Six things the software cannot work out for itself. The first is the one that matters; the rest are smaller and can wait.

QUESTION	WHY IT MATTERS
1 What credit term applies to each of your largest parties? For hospitals: by which day must bills reach them, on which day do they release payment, and is that the same month or the next?	Nothing can be called overdue without it. All ₹8.06 Cr currently sits under "cannot be judged". This is the single input that turns the system on.
2 What are the CRA, CRB, CRC and CRD voucher series? 1,634 rows use them — CRD alone carries ₹1.35 Cr, and 158 of these rows are negative.	They are imported and counted correctly but grouped as "other". Naming them would let receipts and adjustments be told apart from sales.
3 Marg already holds a due date for 19 parties , implying 60-day or 30-day terms. Are those correct and approved?	A yes turns 19 parties judgeable immediately. They are deliberately not applied until a person confirms them.
4 Should a party be chased at 60 days past term, or 90?	The buckets are 60+ today. Changing it is a small edit, but it changes who appears on the call list.
5 Is ₹10,000 the right point below which a balance is not worth chasing? 328 parties sit under it, ₹9.91 L between them.	Those parties are excluded from the call list. The threshold is adjustable in Settings.
6 Who should have a login, and what should each of them be able to change?	Owner, accounts and sales differ in what they may write. Credit terms are restricted deliberately.

One more, worth asking early: for 46 parties Marg's own header total disagrees with the sum of its bill rows, by about ₹1 crore in aggregate. The system follows the header, as Marg does. Does the accounts team know why those differ? It may be nothing, or it may be something worth correcting at source.